

NAIGAMA PARSING SYSTEM (COMPENDIUM)

KEES-JAN HERMANS

This document contains a description of Naigama, a software implementation of a Parsing Expression Grammar compiler, assembler, and bytecode execution engine.



Accompanies release	0.2.1
Author	Kees-Jan Hermans / kees.jan.hermans@gmail.com
Classification	-
Generated on	November 13, 2021

CONTENTS

1. Introduction	5
1.1. Parsing Expression Grammars	5
1.2. Naigama	5
2. Overview	6
2.1. On Grammar	6
2.2. On Assembly	7
2.3. On Bytecode	7
3. Generations	8
4. Grammar	9
4.1. Rules	9
4.2. Expressions	10
4.3. Terms	10
4.4. Matchers	10
5. Assembly	12
5.1. Comments	12
5.2. Labels	12
5.3. Instructions	13
5.4. Parameters	14
6. Bytecode	16
6.1. Bytecode Structure	17
6.2. Opcode Values	17
6.3. Noop Slides and Canaries	18
6.4. Encoding of Parameters	18
7. Output	19
8. Executables	20
8.1. The Compiler	20
8.2. The Assembler	20
8.3. The Engine	20
8.4. The Disassembler	21
9. File Formats	22
9.1. Text Formats	22
9.2. Binary Formats	23
10. Colofon	26
List of Tables	26
References	26
Appendices	27
A. Appendix: Instructions	27
A.1. Instruction: any	29
A.2. Instruction: backcommit	30
A.3. Instruction: call	31
A.4. Instruction: catch	32
A.5. Instruction: char	33
A.6. Instruction: closecapture	34
A.7. Instruction: commit	35

A.8. Instruction: condjump	36
A.9. Instruction: counter	37
A.10. Instruction: end	38
A.11. Instruction: endisolate	39
A.12. Instruction: endreplace	40
A.13. Instruction: fail	41
A.14. Instruction: failtwice	42
A.15. Instruction: intrpcapture	43
A.16. Instruction: isolate	44
A.17. Instruction: jump	45
A.18. Instruction: maskedchar	46
A.19. Instruction: mode	47
A.20. Instruction: noop	48
A.21. Instruction: opencapture	49
A.22. Instruction: partialcommit	50
A.23. Instruction: quad	51
A.24. Instruction: range	52
A.25. Instruction: replace	53
A.26. Instruction: ret	54
A.27. Instruction: scr_add	55
A.28. Instruction: scr_array	56
A.29. Instruction: scr_assign	57
A.30. Instruction: scr_bitand	58
A.31. Instruction: scr_bitnot	59
A.32. Instruction: scr_bitor	60
A.33. Instruction: scr_bitxor	61
A.34. Instruction: scr_builtin	62
A.35. Instruction: scr_call	63
A.36. Instruction: scr_condjump	64
A.37. Instruction: scr_dec	65
A.38. Instruction: scr_div	66
A.39. Instruction: scr_equals	67
A.40. Instruction: scr_gt	68
A.41. Instruction: scr_gteq	69
A.42. Instruction: scr_inc	70
A.43. Instruction: scr_index	71
A.44. Instruction: scr_logand	72
A.45. Instruction: scr_lognot	73
A.46. Instruction: scr_logor	74
A.47. Instruction: scr_lt	75
A.48. Instruction: scr_lteq	76
A.49. Instruction: scr_mul	77
A.50. Instruction: scr_nequals	78
A.51. Instruction: scr_pop	79
A.52. Instruction: scr_pow	80

A.53.	Instruction: scr_push	81
A.54.	Instruction: scr_ret	82
A.55.	Instruction: scr_shift	83
A.56.	Instruction: scr_shiftin	84
A.57.	Instruction: scr_shiftout	85
A.58.	Instruction: scr_string	86
A.59.	Instruction: scr_sub	87
A.60.	Instruction: set	88
A.61.	Instruction: skip	89
A.62.	Instruction: span	90
A.63.	Instruction: testany	91
A.64.	Instruction: testchar	92
A.65.	Instruction: testquad	93
A.66.	Instruction: testset	94
A.67.	Instruction: trap	95
A.68.	Instruction: var	96

1. INTRODUCTION

1.1. **Parsing Expression Grammars.** This is the documentation describing Naigama, which is a collection of specifications, libraries and tools enabling data parsing. It is based, by now loosely, on Lua Parsing Expression Grammars, or LPEG [REF], which implements a Packrat [REF] parsing algorithm.

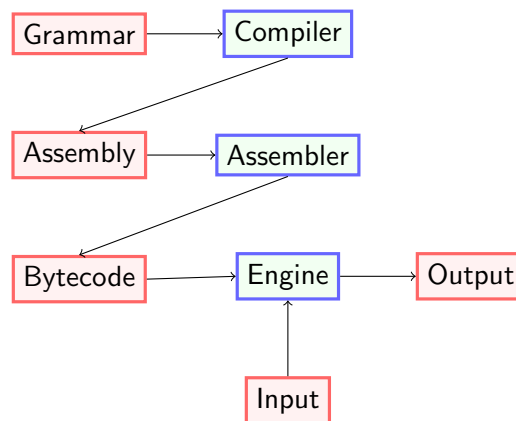
1.2. **Naigama.** Naigama adds to Lua PEG the following features:

- Discoverability features:
 - A fully fledged assembler step in between grammar and bytecode.
 - A capable engine debugger.
- Extra tools:
 - A disassembler.
 - A decompiler.
- Security features:
 - Bit upset event detection and traps.
 - Endless loop detection.
- Interoperability with C (library, header files).
- Added instructions for counters.
- Added instructions for parsing binary inputs.
- A scripting language.

2. OVERVIEW

This document describes the Naigama parsing system, which is a modular system to process structured arbitrary length data inputs, in order to extract meaning (matching, capturing from data), or to manipulate (replacements within data). The system was designed with a focus on information and system security.

The system comes in three parts: a grammar compiler, an assembler, and a bytecode execution engine. Each of the parts has their own input file format specification (grammar description, assembly specification, and bytecode specification), and outputs the format of the next stage. The process in brief: the compiler takes a grammar description that the human end user has made, and turns it into an assembly language. The assembler takes the assembly, and turns it into a bytecode. The engine takes the bytecode and the input, and processes it, resulting either in (various measures of) failure or success.



Schematic overview of Naigama data format transformation and tools

The reason for this modularity is strict separation of tasks, and openness: anyone should be able to take out a module, and replace it with one of their own. It should be expressly possible to take out the compiler, and replace it with another tool that produces the Naigama assembly, for example. Or have a different engine. Or write an optimizer for the assembly.

2.1. On Grammar. Naigama grammar is the human interface of the system. It can be used to define complex syntax definitions in order to match, capture from, and replace in, structured inputs.

People familiar with regular expressions [2], Backus-Naur syntax descriptions [3], and / or Lex and Yacc tools [4], should find this part of the document relatively easy to understand. The ideas underlying Naigama grammar, assembly and bytecode are heavily borrowed from Lua PEG (LPEG) [1] by Roberto Ierusalimsky et al.

2.2. **On Assembly.** The Naigama assembly language is a mixture of two languages: For one, it is the outcome instruction set, according to the aforementioned LPEG whitepaper, of the Naigama grammar compilation process. However, it also contains instructions that won't be produced by this compiler, but may either be produced by a separate optimizer, or other compilers altogether (for example, one that focuses on network packet parsing).

The Naigama assembly language, in all cases, is the human readable variant of the Naigama bytecode, and its instructions are one-on-one translatable from the one to the other, and vice versa (although labels will be lost in the reverse process).

2.3. **On Bytecode.** Naigama bytecode is the machine interface of the system. It runs in the Naigama engine against the input provided by the user. It is designed to be:

- Easy and quick to interpret.
- Resilient against bit upset events.
- Resilient against endless loops.
- Usable while keeping both bytecode and input as read-only buffers.

3. GENERATIONS

Naigama is written in generations. That is to say: each subsequent generation of Naigama uses the tools generated in the previous one, to build itself. Initially, there's a bootstrap compiler, followed by compilers that use their own grammar to parse their own grammar, followed by one that can completely compile itself.

This leads to the following sequence:

- Generation zero: This consists solely of a grammar compiler and assembler written in perl.
- Generation one: This generation is the grammar only tooling built in C, based on the bytecode to parse grammar grammar and assembly grammar, generated by generation zero.
- Generation two: equal to generation one, but then based on the bytecode generated by the tooling in generation one.
- Generation three: native tooling, including scripting, based on the bytecode generated by the generation two tooling.

4. GRAMMAR

A Naigama grammar file or buffer consists either of a single unnamed expression, or of a sequence of named expressions (called rules), whitespace and comments, denoted in ASCII text. Naigama grammar is defined with the express purpose of matching, validating and/or capturing from an arbitrary data input. Naigama grammar is taken by the Naigama compiler program (naic) and turned into Naigama assembly [section 5].

```
-- JSON; JavaScript Object Notation.
TOP          <- JSON
__prefix     <- %s*
JSON         <- HASH END
HASH         <- CBOPEN OPTHASHELTS CBCLOSE
OPTHASHELTS  <- HASHELTS / ...
HASHELTS     <- HASHELT COMMA HASHELTS / HASHELT
HASHELT      <- STRING COLON VALUE
ARRAY        <- ABOPEN OPTARRAYELTS ABCLOSE
OPTARRAYELTS <- ARRAYELTS / ...
ARRAYELTS    <- VALUE COMMA ARRAYELTS / VALUE
VALUE        <- STRING / FLOAT / INT / BOOL / NULL / HASH / ARRAY
STRING       <- { '"' ( '\\' ([nrtv"] / [0-9]^3) / [^"\\] ) * '"' }
INT          <- { [0-9]+ }
FLOAT        <- { [0-9]* '.' [0-9]+ }
BOOL         <- { 'true' / 'false' }
NULL         <- { 'null' }
CBOPEN       <- '{'
CBCLOSE      <- '}'
ABOPEN       <- '['
ABCLOSE      <- ']'
COMMA        <- ','
COLON        <- ':'
END          <- !.
```

Example of a Naigama grammar meant to validate and capture from JSON[6] text.

4.1. Rules. A rule is defined as an *Identifier*, followed by a left-pointing arrow (composed of a less-than and a minus sign), followed by a matching expression.

When a Naigama grammar consists of a sequence of rules (as opposed to a single line expression), the first rule is used as the starting point for matching inputs.

```
RULE1 <- 'a' / RULE2
RULE2 <- 'b' / RULE3
```

```
RULE3 <- 'c'
```

Definition of a set of rules

4.1.1. *Identifiers.* Identifiers, in Naigama, are defined as a combination of letters, numbers and the underscore character, not starting with a number, of between one and 64 characters long.

```
IDENTIFIER <- [a-zA-Z_][a-zA-Z0-9_]{1,63}
```

Definition of an identifier

Identifiers are used to start rule definitions, and as references to rules in expressions.

4.1.2. *Special Rules.* There is currently one special rule in Naigama grammar: a rule called '`__prefix`' will be treated specially. This rule will be, from its definition onwards, called before the execution of any subsequent rule definition expression. The aim is to make whitespace and comment filtering in program language parsing easier. See also the earlier 'JSON' grammar example.

4.2. **Expressions.** Expressions are lists of terms, optionally separated by the OR operator (denoted by the forward slash sign).

4.3. **Terms.**

4.4. **Matchers.**

4.4.1. *NOT and AND Modifiers.*

4.4.2. *Quantifiers.*

4.4.3. *The ANY Matcher.* The ANY matcher also allows you to define end-of-input, like so:

```
ENDOFINPUT <- !.
```

Definition of end-of-input

4.4.4. *The SET Matcher.*

4.4.5. *The STRING Matcher.*

4.4.6. *The BITMASK Matcher.*

4.4.7. *The HEXLITERAL Matcher.*

4.4.8. *The VARCAPTURE Matcher.*

4.4.9. *The CAPTURE Matcher.*

4.4.10. *The GROUP Matcher.*

4.4.11. *The MACRO Matcher.*

4.4.12. *The VARREFERENCE Matcher.*

4.4.13. *The REFERENCE Matcher.*

5. ASSEMBLY

A Naigama assembly file or buffer consists of sequences of whitespace, comments, labels, and (parametrized) instructions, separated by new lines, and denoted in ASCII text.

Naigama assembly is taken by the Naigama assembler program (naia) and turned into bytecode.

```
-- Compilation of: TEST <- { 'a' } { 'a' } { 'a' / 'b' }
  call TEST
  end
-- Rule
TEST:
  opencapture 0
  char 61
  closecapture 0 0
  opencapture 1
  char 61
  closecapture 1 0
  opencapture 2
  catch __LABEL_30 -- alternative
  char 61
  commit __LABEL_31
__LABEL_30:
  char 62
__LABEL_31:
  closecapture 2 0
  ret
```

Example of a piece of Naigama assembly

5.1. Comments. Just like in Naigama grammar, comments in Naigama assembly start with two minus signs, and end at a new line. They can be given on separate lines, or they can be postfixed to instructions or labels.

5.2. Labels. Labels are identifiers (the same identifiers as in their definition in the grammar chapter) or numbers, followed by a colon sign. They are used as positions for the bytecode to jump to, when used by instructions.

When performing the assembly, the Naigama assembler program resolves each label to their offset in the resultant bytecode, and replaces each reference to a label with that offset.

Labels don't quite disappear though - you have the option to make the assembler program emit a so called 'labelmap' file, which stores the old label names, mapped to the bytecode offsets given by the assembler, and which can be used later for debugging purposes by the bytecode execution engine.

When disassembling (bytecode to assembly), each instruction is always prefixed by a label in the form of the instruction's offset in decimal, so that jumps are always correct (but non descriptive).

5.3. Instructions.

Naigama Assembly Instructions

Mnemonic	Param1	Param2
'any'		
'backcommit'		LABEL
'call'		LABEL
'catch'		LABEL
'char'	char	
'closecapture'	slot	
'commit'		LABEL
'condjump'	register	LABEL
'counter'	register	value
'end'	code	
'endisolate'		
'endreplace'		
'fail'		
'failtwice'		
'intrpcapture'		
'isolate'	slot	
'jump'		LABEL
'maskedchar'	char	mask
'mode'		
'noop'		
'opencapture'	slot	
'partialcommit'		LABEL
'quad'	quad	
'range'	from	until
'replace'	slot	LABEL
'ret'		
'scr_add'		
'scr_array'		
'scr_assign'		
'scr_bitand'		
'scr_bitnot'		

Mnemonic	Param1	Param2
'scr_bitor'		
'scr_bitxor'		
'scr_builtin'		
'scr_call'		
'scr_condjump'		
'scr_dec'		
'scr_div'		
'scr_equals'		
'scr_gt'		
'scr_gteq'		
'scr_inc'		
'scr_index'		
'scr_logand'		
'scr_lognot'		
'scr_logor'		
'scr_lt'		
'scr_lteq'		
'scr_mul'		
'scr_nequals'		
'scr_pop'		
'scr_pow'		
'scr_push'		
'scr_ret'		
'scr_shift'		
'scr_shiftin'		
'scr_shiftout'		
'scr_string'		
'scr_sub'		
'set'	set	
'skip'	number	
'span'	set	
'testany'		LABEL
'testchar'	char	LABEL
'testquad'	quad	LABEL
'testset'	set	LABEL
'trap'		
'var'	slot	

See table [A].

5.4. **Parameters.** Parameters follow the instruction in assembly text. They are separated from the instruction and any other parameters by spaces. The following types of parameters exist:

5.4.1. *Characters*. Character parameters are denoted as two byte hexadecimal values.

5.4.2. *Quads*. Quad parameters are denoted as eight byte hexadecimal values.

5.4.3. *Sets*. Set parameters are denoted as 64 byte hexadecimal values, representing a bit-mask of 256 possible booleans.

5.4.4. *Registers, Slots, Codes and Numbers*. These parameters are all denoted as decimal numbers.

6. BYTECODE

Naigama Bytecode Instructions

Mnemonic	Opcode	Param1	Param2	Length
any	000003e4			4
backcommit	000403c0	address		8
call	00040382	address		8
catch	00040393	address		8
char	000403d7	char		8
closecapture	00040300	slot		8
commit	00040336	address		8
condjump	00080321	register	address	12
counter	00080356	register	value	12
end	000400d8	code		8
endisolate	00003005			4
endreplace	00000399			4
fail	0000034b			4
failtwice	00000390			4
intrpcapture	0008000f			12
isolate	00043003	slot		8
jump	00040333	address		8
maskedchar	00080365	char	mask	12
mode	0004000a			8
noop	00000000			4
opencapture	0004039c	slot		8
partialcommit	000403b4	address		8
quad	0004037e	quad		8
range	000803bd	from	until	12
replace	00080348	slot	address	12
ret	000003a0			4
scr_add	0000050c			4
scr_array	00040006			8
scr_assign	000005c9			4
scr_bitand	00000527			4
scr_bitnot	00000574			4
scr_bitor	0000053c			4
scr_bitxor	0000052d			4
scr_builtin	000407cf			8
scr_call	00040503			8
scr_condjump	0004000c			8
scr_dec	0000053a			4
scr_div	00000581			4
scr_equals	0000056c			4
scr_gt	0000056f			4

Mnemonic	Opcode	Param1	Param2	Length
scr_gteq	0000054d			4
scr_inc	000005f6			4
scr_index	00000009			4
scr_logand	0000052e			4
scr_lognot	000005f9			4
scr_logor	000005a9			4
scr_lt	00000595			4
scr_lteq	00000522			4
scr_mul	0000058b			4
scr_nequals	00000572			4
scr_pop	000005cc			4
scr_pow	00000542			4
scr_push	000c0003			16
scr_ret	00000555			4
scr_shift	00040005			8
scr_shiftin	0000057d			4
scr_shiftout	00000517			4
scr_string	000017bb			4
scr_sub	000005bb			4
set	002003ca	set		36
skip	00040330	number		8
span	002003e1	set		36
testany	00040306	address		8
testchar	0008039a	address	char	12
testquad	000803db	address	quad	12
testset	00240363	address	set	40
trap	ff00ffff			4
var	000403ee	slot		8

6.1. Bytecode Structure. A Naigama bytecode file or buffer consists of a sequence of binary instructions which, in turn, each consist of a binary encoded opcode, plus their parameters, should they have any.

The amount and kind of parameters following an opcode, is strictly defined: the same opcode will always be followed by the same kinds of parameters and therefore, an instruction type will always be the same size (see table [6]).

Naigama bytecode is taken by the Naigama engine program (naie) or library, and run against an input, to produce an output.

6.2. Opcode Values. Opcode values are determined through

- Grouping;
- Hamming distance;

- Instruction size;

6.3. **Noop Slides and Canaries.** Implementations that want each instruction to have exactly the same size, can choose to pad the encoding of shorter instructions with either no-ops, or canaries.

6.4. **Encoding of Parameters.**

6.4.1. *Address.*

6.4.2. *Char.*

6.4.3. *Slot.*

6.4.4. *Register.*

7. OUTPUT

8. EXECUTABLES

8.1. **The Compiler.** The Naigama grammar compiler is called 'naic'. It takes a grammar file as input, and outputs assembly text. It can be invoked as follows:

```
$ naic -i myfile.niag -o myfile.asm
```

Example of the invocation of the Naigama grammar compiler

Bear in mind that both the input (grammar) file, as well as the output (assembly) file may be omitted (in which case they will be assumed to be stdin and stdout, respectively), or be denoted as a minus sign ('-'), which will have the same effect.

You may also use the following options:

- -m <path> Tells the compiler to emit the slotmap file in <path>.
- -D Creates a lot of debugging output.
- -t Tells the compiler to surround generated rule code with 'trap' instructions.

8.2. **The Assembler.** The Naigama grammar assembler is called 'naia'. It takes an assembly file as input, and outputs bytecode. It can be invoked as follows:

```
$ naia -i myfile.asm -o myfile.byc
```

Example of the invocation of the Naigama assembler

Bear in mind that both the input (assembly) file, as well as the output (bytecode) file may be omitted (in which case they will be assumed to be stdin and stdout, respectively), or be denoted as a minus sign ('-'), which will have the same effect.

You may also use the following options:

- -l <path> Tells the assembler to emit the labelmap file in <path>.
- -D Creates a lot of debugging output.

8.3. **The Engine.** The Naigama bytecode execution engine is called 'naie'. It takes a bytecode file as input, as well as an input file. It can be invoked as follows:

```
$ naive -c myfile.byc -i myfile.dat -o myfile.out
```

Example of the invocation of the Naigama engine

8.4. **The Disassembler.** The Naigama disassembler is called 'naid'. It takes a bytecode file as input, and outputs assembly.

```
$ naid -i myfile.byc -o myfile.asm
```

Example of the invocation of the Naigama disassembler

The output of the disassembler will differ from the assembly generated by the compiler in that:

- Textual labels will be gone, instead:
- Every instruction will be prefixed by a numeric label which is identical to that instruction's position offset in the bytecode, and
- All jumps will therefore also be using those position labels.

9. FILE FORMATS

This section describes all the user available memory and file formats associated with Naigama.

9.1. Text Formats.

9.1.1. *Grammar*. The Naigama grammar text format is descibed in [REF]. Various pre-compilers can potentially produce these format however, these are out of scope of this documentation. The grammar format is consumed by the Naigama assembler, naic. For the use of naic, see [REF].

9.1.2. *Assembly*. The Naigama assembly text format is descibed in [REF]. This format is produced by the Naigama compiler, naic, and is consumed by the Naigama assembler, naia. For the use of naia, see [REF].

9.1.3. *Slotmap.h*. The Naigama compiler assigns a slot number to each capture it encounters. It then assigns an internal unique name to those capture numbers, or slots. Finally, the compiler can be requested, through the command line, to 'rain down' these named mappings to slot number in a C-header file, so that API users can use easy-to-remember defines to address their captures.

9.1.4. *Disassembly*. The Naigama disassembler, naid, will take bytecode and reproduce the assembly. This assembly will be different from the original assembly in that:

- It will not contain textually understandable labels.
- It will prefix every instruction with a numeric label, which is equivalent to the byte-code offset of the instruction's opcode.
- These offsets will then also be used in jumps.

9.1.5. *Engine Debug Logs*. Running the Naigama engine, naie, in debug mode will yield, on standard error, a log, which will contain, per instruction executed, a line representing the engine's internal state, most notably its stack, like so:

```
CHAR bc 2316 in 482 0403020100\_\_\_\_\_ st (014 prec.) ALT:1484
CLL:1476 C LL:1820 ALT:1944 CLL:1928 ALT:1652 CLL:1644 CLL:2440
```

These lines are built up as follows:

- The instruction being executed (in this case, 'char').
- The bytecode offset (in this case, 2316 decimal).
- The input offset (in this case, 482 decimal).

- A subset of the input, from the input offset (in this in hexadecimal).
- An overview of the stack (in this case, with 14 preceding items left out), followed by the last stack items, either 'call' items (with their return addresses), or 'alternative' items (with their jump-to addresses in case they catch a FAIL condition).

9.2. Binary Formats. This section describes the formats used by the Naigama tooling, other than the ones described in the sections on grammar, assembly and bytecode.

9.2.1. Slotmap Format. The slotmap file exists to help developers by easing access to captures in complex grammars.

The Naigama grammar compiler can be instructed to emit a slotmap file. This is a file which maps a unique name to a capture region's index. This is provided so that, instead of using the index of capture region (which requires hand counting them in your grammar file, which can be tedious and error-prone, and something that would not survive a grammar reshuffle, or the introduction of a capture region before the one you're interested in), you can use a naming scheme for your capture regions.

Names in the slotmap file are bound semantically to the capture region: they are made up of the name of the rule, an underscore, and all alphabetic characters in the capture region, cast to uppercase. Should any name occur twice, it will be postfixed with a counter until it's unique.

For example, the following rule with capture region definition:

```
RULE <- { IDENT } OPTARGS LEFTARROW EXPRESSION
```

will result in slotmap identifier 'RULE_IDENT'.

The binary format of a slotmap file is composed as a sequence of binary records: the slot index, denoted as a 32 bit network order unsigned integer, a field of 32 bit all ones, and the slot name, denoted as a zero-terminated string.

```
00 00 00 00 ff ff ff ff 52 55 4c 45 5f 49 44 45 .....RULE_IDE
4e 54 00 00 00 00 01 ff ff ff ff 45 58 50 52 45 NT.....EXPRE
53 53 49 4f 4e 5f 54 45 52 4d 53 00 00 00 00 02 SSION_TERMS.....
ff ff ff ff 45 58 50 52 45 53 53 49 4f 4e 5f 54 ....EXPRESSION_T
45 52 4d 53 5f 31 00 00 00 00 03 ff ff ff ff 45 ERMS_1.....E
58 50 52 45 53 53 49 4f 4e 5f 54 45 52 4d 53 5f XPRESSION_TERMS_
32 00 00 00 00 04 ff ff ff ff 54 45 52 4d 53 5f 2.....TERMS_
54 45 52 4d 00 00 00 00 05 ff ff ff ff 54 45 52 TERM.....TER
4d 5f 4e 4f 54 41 4e 44 00 00 00 00 06 ff ff ff M_NOTAND.....
```

Example of the head of a slotmap file, hex dumped

9.2.2. *Labelmap Format.* The labelmap format is optionally emitted by the assembler, and exists to:

- Allow for better debugging, because offsets in bytecode can be reduced to more intuitively named sections (rule names always make it to the labelmap unchanged).
- Allow for calling the bytecode at symbols directly, which in turn allows you to use a bytecode blob more as a database of functions.

The binary format of the labelmap file is composed as a sequence of binary records: the bytecode offset, denoted as a 32 bit network order unsigned integer, followed by the labelstring, followed by a zero byte.

```

00 00 00 10 47 52 41 4d 4d 41 52 00 00 00 00 28 ....GRAMMAR....(
5f 5f 4c 41 42 45 4c 5f 31 31 00 00 00 00 38 5f __LABEL_11....8_
5f 4c 41 42 45 4c 5f 31 32 00 00 00 00 40 5f 5f _LABEL_12....@__
4c 41 42 45 4c 5f 36 00 00 00 00 48 5f 5f 4c 41 LABEL_6....H__LA
42 45 4c 5f 37 00 00 00 00 54 5f 5f 70 72 65 66 BEL_7....T__pref
69 78 00 00 00 00 5c 5f 5f 4c 41 42 45 4c 5f 32 ix....\__LABEL_2
39 00 00 00 00 7c 5f 5f 4c 41 42 45 4c 5f 34 33 9....|__LABEL_43
00 00 00 00 a8 5f 5f 4c 41 42 45 4c 5f 34 34 00 .....__LABEL_44.
00 00 00 b8 5f 5f 4c 41 42 45 4c 5f 33 32 00 00 .....__LABEL_32..
00 00 e4 5f 5f 4c 41 42 45 4c 5f 35 35 00 00 00 .....__LABEL_55...
01 10 5f 5f 4c 41 42 45 4c 5f 35 36 00 00 00 01 ..__LABEL_56....
10 5f 5f 4c 41 42 45 4c 5f 33 33 00 00 00 01 18 __LABEL_33.....
5f 5f 4c 41 42 45 4c 5f 33 30 00 00 00 01 1c 45 __LABEL_30.....E
4e 44 00 00 00 01 34 5f 5f 4c 41 42 45 4c 5f 35 ND....4__LABEL_5
38 00 00 00 01 38 44 45 46 49 4e 49 54 49 4f 4e 8....8DEFINITION
00 00 00 01 4c 53 49 4e 47 4c 45 5f 45 58 50 52 ....LSINGLE_EXPR
45 53 53 49 4f 4e 00 00 00 01 60 52 55 4c 45 00 ESSION....'RULE.
00 00 01 a0 45 58 50 52 45 53 53 49 4f 4e 00 00 ....EXPRESSION..
00 01 e4 5f 5f 4c 41 42 45 4c 5f 31 30 36 00 00 ...__LABEL_106..

```

Example of a section of labelmap file, hex dumped

9.2.3. *Engine Output Format.* The Naigama bytecode execution engine, naie, produces, on success, output for digital processing, in the form of a binary table, which is structured as follows:

- Each record is four 32 bit integers, in network order.
- The first record contains the end code of the matching process. This is the same code as was given as a parameter to the 'end' instruction that resulted in the execution finishing. This is the first field, by default it is zero. The second field contains the amount of subsequent records.
- Subsequent records have as fields, either:

- Type, slot, start, length (when of 'capture' type), or:
- Type, slot, start, char (when of 'replace' type).

Bear in mind the following:

- 'Capture' type is denoted as 1, 'Replace' as 3.
- Offsets and lengths of captures refer to the input buffer.
- Offsets and lengths of replacements refer to the bytecode.

```

00 00 00 00 00 00 03 64 00 00 00 00 00 00 00 00 .....d.....
00 00 00 01 00 00 00 00 00 00 00 00 1c 00 00 00 23 .....#
00 00 00 01 00 00 00 03 00 00 00 32 00 00 00 59 .....2...Y
00 00 00 01 00 00 00 04 00 00 00 32 00 00 00 55 .....2...U
00 00 00 01 00 00 00 14 00 00 00 32 00 00 00 55 .....2...U
00 00 00 01 00 00 00 02 00 00 00 34 00 00 00 3f .....4...?
00 00 00 01 00 00 00 04 00 00 00 34 00 00 00 3f .....4...?
00 00 00 01 00 00 00 17 00 00 00 34 00 00 00 3e .....4...>
00 00 00 01 00 00 00 06 00 00 00 3e 00 00 00 3f .....>...?
00 00 00 01 00 00 00 03 00 00 00 42 00 00 00 53 .....B...S
00 00 00 01 00 00 00 04 00 00 00 42 00 00 00 53 .....B...S
00 00 00 01 00 00 00 17 00 00 00 42 00 00 00 53 .....B...S

```

Example of the head of engine output with a zero end code and 868 matches (all captures), hex dumped

10. COLOFON

LIST OF TABLES

REFERENCES

- [1] A Text Pattern-Matching Tool based on Parsing Expression Grammars <https://www.inf.puc-rio.br/~roberto/docs/peg.pdf>
- [2] Regular Expressions https://en.wikipedia.org/wiki/Regular_expression
- [3] Backus Naur Form https://en.wikipedia.org/wiki/Backus-Naur_form
- [4] Yacc Yet Another Compiler Compiler <https://en.wikipedia.org/wiki/Yacc>
- [5] JavaScript, or ECMAScript <https://www.ecma-international.org/publications/standards/Ecma-262.htm>
- [6] JSON, JavaScript Object Notation <https://www.json.org/>
- [7] Perl, the Perl Programming Language <https://www.perl.org/>

Appendices

A. APPENDIX: INSTRUCTIONS

Naigama Assembly Instructions

Mnemonic	Param1	Param2
'any'		
'backcommit'		LABEL
'call'		LABEL
'catch'		LABEL
'char'	char	
'closecapture'	slot	
'commit'		LABEL
'condjump'	register	LABEL
'counter'	register	value
'end'	code	
'endisolate'		
'endreplace'		
'fail'		
'failtwice'		
'intrpcapture'		
'isolate'	slot	
'jump'		LABEL
'maskedchar'	char	mask
'mode'		
'noop'		
'opencapture'	slot	
'partialcommit'		LABEL
'quad'	quad	
'range'	from	until
'replace'	slot	LABEL
'ret'		
'scr_add'		
'scr_array'		
'scr_assign'		
'scr_bitand'		
'scr_bitnot'		
'scr_bitor'		
'scr_bitxor'		
'scr_builtin'		
'scr_call'		
'scr_condjump'		
'scr_dec'		

Mnemonic	Param1	Param2
'scr_div'		
'scr_equals'		
'scr_gt'		
'scr_gteq'		
'scr_inc'		
'scr_index'		
'scr_logand'		
'scr_lognot'		
'scr_logor'		
'scr_lt'		
'scr_lteq'		
'scr_mul'		
'scr_nequals'		
'scr_pop'		
'scr_pow'		
'scr_push'		
'scr_ret'		
'scr_shift'		
'scr_shiftin'		
'scr_shiftout'		
'scr_string'		
'scr_sub'		
'set'	set	
'skip'	number	
'span'	set	
'testany'		LABEL
'testchar'	char	LABEL
'testquad'	quad	LABEL
'testset'	set	LABEL
'trap'		
'var'	slot	

A.1. Instruction: **any**.

A.1.1. *Mode*. This instruction is available in mode 0 (parser).

A.1.2. *Summary*. This instruction advances the input pointer by one, if possible. If successful, the input pointer is advanced and the next instruction is executed. If not successful (because end-of-input has not been reached) then the FAIL state is invoked.

A.1.3. *Grammar and Compiling*. The 'any' instruction is emitted whenever the 'dot' (.) token is used in rule production.

A.1.4. *Assembly Syntax*.

```
ANYINSTR <- { 'any' }
```

A.1.5. *Bytecode Encoding*. This instruction has a size of 4 bytes and is structured in bytecode as follows:

00	00	00	03	e4
----	----	----	----	----

A.1.6. *Execution State Change*. .

Original state: (p, i, e, c)

Operation: **any** ; $i < |S|$

Failure state: **(Fail, i, e, c)**

Success state: $(p + 1, i + 1, e, c)$

A.2. Instruction: **backcommit**.

A.2.1. *Mode*. This instruction is available in mode 0 (parser).

A.2.2. *Summary*.

A.2.3. *Grammar and Compiling*.

A.2.4. *Assembly Syntax*.

```
BACKCOMMITINSTR <- { 'backcommit' } S LABEL
S <- %s+
LABEL <- { [a-zA-Z0-9_]^1-64 }
```

A.2.5. *Bytecode Encoding*. This instruction has a size of 8 bytes and is structured in bytecode as follows:

00	00	04	03	c0
04	00	00	00	00

A.2.6. *Execution State Change*. .

A.3. Instruction: **call**.

A.3.1. *Mode*. This instruction is available in mode 0 (parser).

A.3.2. *Summary*.

A.3.3. *Grammar and Compiling*.

A.3.4. *Assembly Syntax*.

```
CALLINSTR <- { 'call' } S LABEL
S <- %s+
LABEL <- { [a-zA-Z0-9_]^1-64 }
```

A.3.5. *Bytecode Encoding*. This instruction has a size of 8 bytes and is structured in bytecode as follows:

00	00	04	03	82
04	00	00	00	00

A.3.6. *Execution State Change*. .

A.4. Instruction: **catch**.

A.4.1. *Mode*. This instruction is available in mode 0 (parser).

A.4.2. *Summary*.

A.4.3. *Grammar and Compiling*.

A.4.4. *Assembly Syntax*.

```
CATCHINSTR <- { 'catch' } S LABEL
S <- %s+
LABEL <- { [a-zA-Z0-9_]^1-64 }
```

A.4.5. *Bytecode Encoding*. This instruction has a size of 8 bytes and is structured in bytecode as follows:

00	00	04	03	93
04	00	00	00	00

A.4.6. *Execution State Change*. .

A.5. Instruction: **char**.

A.5.1. *Mode*. This instruction is available in mode 0 (parser).

A.5.2. *Summary*.

A.5.3. *Grammar and Compiling*.

A.5.4. *Assembly Syntax*.

```
CHARINSTR <- { 'char' } S HEXBYTE
S <- %s+
HEXBYTE <- { [0-9a-fA-F]^2 }
```

A.5.5. *Bytecode Encoding*. This instruction has a size of 8 bytes and is structured in bytecode as follows:

00	00	04	03	d7
04	00	00	00	00

A.5.6. *Execution State Change*. .

A.6. Instruction: **closecapture**.

A.6.1. *Mode*. This instruction is available in mode 0 (parser).


A.6.2. *Summary*.

A.6.3. *Grammar and Compiling*.

A.6.4. *Assembly Syntax*.

```
CLOSECAPTUREINSTR <- { 'closecapture' } S SLOT
S <- %s+
SLOT <- UNSIGNED
```

A.6.5. *Bytecode Encoding*. This instruction has a size of 8 bytes and is structured in bytecode as follows:



A.6.6. *Execution State Change*. .

A.7. Instruction: **commit**.

A.7.1. *Mode*. This instruction is available in mode 0 (parser).

A.7.2. *Summary*.

A.7.3. *Grammar and Compiling*.

A.7.4. *Assembly Syntax*.

```
COMMITINSTR <- { 'commit' } S LABEL
S <- %s+
LABEL <- { [a-zA-Z0-9_]^1-64 }
```

A.7.5. *Bytecode Encoding*. This instruction has a size of 8 bytes and is structured in bytecode as follows:

00	00	04	03	36
04	00	00	00	00

A.7.6. *Execution State Change*. .

A.8. Instruction: **condjump**.

A.8.1. *Summary*.

A.8.2. *Grammar and Compiling*.

A.8.3. *Assembly Syntax*.

```
CONDJUMPINSTR <- { 'condjump' } S REGISTER S LABEL
S <- %s+
REGISTER <- UNSIGNED
LABEL <- { [a-zA-Z0-9_]^1-64 }
```

A.8.4. *Bytecode Encoding*. This instruction has a size of 12 bytes and is structured in bytecode as follows:

00	00	08	03	21
04	00	00	00	00
08	00	00	00	00

A.8.5. *Execution State Change*. .

A.9. Instruction: counter.

A.9.1. *Mode.* This instruction is available in mode 0 (parser).

A.9.2. *Summary.*

A.9.3. *Grammar and Compiling.*

A.9.4. *Assembly Syntax.*

```
COUNTERINSTR <- { 'counter' } S REGISTER S UNSIGNED
S <- %s+
REGISTER <- UNSIGNED
UNSIGNED <- { [0-9]+ }
```

A.9.5. *Bytecode Encoding.* This instruction has a size of 12 bytes and is structured in bytecode as follows:

00	00	08	03	56
04	00	00	00	00
08	00	00	00	00

A.9.6. *Execution State Change.* .

A.10. Instruction: end.

A.10.1. *Mode.* This instruction is available in both mode 0 (parser) and mode 1 (script interpreter).


A.10.2. *Summary.*

A.10.3. *Grammar and Compiling.*

A.10.4. *Assembly Syntax.*

```
ENDINSTR <- { 'end' } S CODE
S <- %s+
CODE <- UNSIGNED
```

A.10.5. *Bytecode Encoding.* This instruction has a size of 8 bytes and is structured in bytecode as follows:



A.10.6. *Execution State Change.* .

A.11. **Instruction: endisolate.**

A.11.1. *Mode.* This instruction is available in mode 0 (parser).

A.11.2. *Summary.*

A.11.3. *Grammar and Compiling.*

A.11.4. *Assembly Syntax.*

```
ENDISOLATEINSTR <- { 'endisolate' }
```

A.11.5. *Bytecode Encoding.* This instruction has a size of 4 bytes and is structured in bytecode as follows:

00	00	00	30	05
----	----	----	----	----

A.11.6. *Execution State Change.* .

A.12. **Instruction: endreplace.**

A.12.1. *Mode.* This instruction is available in mode 0 (parser).

A.12.2. *Summary.*

A.12.3. *Grammar and Compiling.*

A.12.4. *Assembly Syntax.*

```
ENDREPLACEINSTR <- { 'endreplace' }
```

A.12.5. *Bytecode Encoding.* This instruction has a size of 4 bytes and is structured in bytecode as follows:

00

00

00

03

99

A.12.6. *Execution State Change.* .

A.13. **Instruction: fail.**

A.13.1. *Mode.* This instruction is available in mode 0 (parser).

A.13.2. *Summary.*

A.13.3. *Grammar and Compiling.*

A.13.4. *Assembly Syntax.*

```
FAILINSTR <- { 'fail' }
```

A.13.5. *Bytecode Encoding.* This instruction has a size of 4 bytes and is structured in bytecode as follows:

00	00	00	03	4b
----	----	----	----	----

A.13.6. *Execution State Change.* .

A.14. **Instruction: failtwice.**

A.14.1. *Mode.* This instruction is available in mode 0 (parser).

A.14.2. *Summary.*

A.14.3. *Grammar and Compiling.*

A.14.4. *Assembly Syntax.*

```
FAILTWICEINSTR <- { 'failtwice' }
```

A.14.5. *Bytecode Encoding.* This instruction has a size of 4 bytes and is structured in bytecode as follows:

00	00	00	03	90
----	----	----	----	----

A.14.6. *Execution State Change.* .

A.15. Instruction: `intrpcapture`.

A.15.1. *Summary*.

A.15.2. *Grammar and Compiling*.

A.15.3. *Assembly Syntax*.

```
INTRPCAPTUREINSTR <- { 'intrpcapture' } S INTRPCAPTURETYPES
S <- %s+
INTRPCAPTURETYPES <- { 'ruint32' }
```

A.15.4. *Bytecode Encoding*. This instruction has a size of 12 bytes and is structured in bytecode as follows:

00	00	08	00	0f
04	00	00	00	00
08	00	00	00	00

A.15.5. *Execution State Change*. .

A.16. Instruction: isolate.

A.16.1. *Mode*. This instruction is available in mode 0 (parser).


A.16.2. *Summary*.

A.16.3. *Grammar and Compiling*.

A.16.4. *Assembly Syntax*.

```
ISOLATEINSTR <- { 'isolate' } S SLOT
S <- %s+
SLOT <- UNSIGNED
```

A.16.5. *Bytecode Encoding*. This instruction has a size of 8 bytes and is structured in bytecode as follows:



A.16.6. *Execution State Change*. .

A.17. Instruction: **jump**.

A.17.1. *Mode*. This instruction is available in both mode 0 (parser) and mode 1 (script interpreter).

A.17.2. *Summary*.

A.17.3. *Grammar and Compiling*.

A.17.4. *Assembly Syntax*.

```
JUMPINSTR <- { 'jump' } S LABEL
S <- %s+
LABEL <- { [a-zA-Z0-9_]^1-64 }
```

A.17.5. *Bytecode Encoding*. This instruction has a size of 8 bytes and is structured in bytecode as follows:

00	00	04	03	33
04	00	00	00	00

A.17.6. *Execution State Change*. .

A.18. Instruction: **maskedchar**.

A.18.1. *Mode*. This instruction is available in mode 0 (parser).

A.18.2. *Summary*.

A.18.3. *Grammar and Compiling*.

A.18.4. *Assembly Syntax*.

```
MASKEDCHARINSTR <- { 'maskedchar' } S HEXBYTE S HEXBYTE
S <- %s+
HEXBYTE <- { [0-9a-fA-F]^2 }
```

A.18.5. *Bytecode Encoding*. This instruction has a size of 12 bytes and is structured in bytecode as follows:

00	00	08	03	65
04	00	00	00	00
08	00	00	00	00

A.18.6. *Execution State Change*. .

A.19. Instruction: mode.

A.19.1. *Mode*. This instruction is available in both mode 0 (parser) and mode 1 (script interpreter).

A.19.2. *Summary*. This instruction changes the mode of the bytecode interpreter to either parser mode (mode 0) or script interpreter mode (mode 1). Different modes have mostly different instruction sets, although they share a few instructions (noop, end, and mode).

A.19.3. *Grammar and Compiling*. This instruction is emitted when the compiler goes from rule compilation (mode 0) to script compilation (mode 1) sections, and vice versa.

A.19.4. *Assembly Syntax*.

```
MODEINSTR <- { 'mode' } S NUMBER
S <- %s+
NUMBER <- { '-'? [0-9]+ }
```

A.19.5. *Bytecode Encoding*. This instruction has a size of 8 bytes and is structured in bytecode as follows:

00	00	04	00	0a
04	00	00	00	00

A.19.6. *Execution State Change*. .

A.20. Instruction: **noop**.

A.20.1. *Mode*. This instruction is available in both mode 0 (parser) and mode 1 (script interpreter).

A.20.2. *Summary*. This instruction does nothing.

A.20.3. *Grammar and Compiling*. This instruction may be emitted to align bytecode.

A.20.4. *Assembly Syntax*.

```
NOOPINSTR <- { 'noop' }
```

A.20.5. *Bytecode Encoding*. This instruction has a size of 4 bytes and is structured in bytecode as follows:

00

00

00

00

00

A.20.6. *Execution State Change*. .

Original state: (p, i, e, c)

Resultant state: (p, i, e, c)

A.21. Instruction: **opencapture**.

A.21.1. *Mode*. This instruction is available in mode 0 (parser).

A.21.2. *Summary*.

A.21.3. *Grammar and Compiling*.

A.21.4. *Assembly Syntax*.

```
OPENCAPTUREINSTR <- { 'opencapture' } S SLOT
S <- %s+
SLOT <- UNSIGNED
```

A.21.5. *Bytecode Encoding*. This instruction has a size of 8 bytes and is structured in bytecode as follows:

00	00	04	03	9c
04	00	00	00	00

A.21.6. *Execution State Change*. .

A.22. Instruction: **partialcommit**.

A.22.1. *Mode*. This instruction is available in mode 0 (parser).


A.22.2. *Summary*.

A.22.3. *Grammar and Compiling*.

A.22.4. *Assembly Syntax*.

```
PARTIALCOMMITINSTR <- { 'partialcommit' } S LABEL
S <- %s+
LABEL <- { [a-zA-Z0-9_]^1-64 }
```

A.22.5. *Bytecode Encoding*. This instruction has a size of 8 bytes and is structured in bytecode as follows:



A.22.6. *Execution State Change*. .

A.23. Instruction: quad.

A.23.1. *Mode*. This instruction is available in mode 0 (parser).

A.23.2. *Summary*.

A.23.3. *Grammar and Compiling*.

A.23.4. *Assembly Syntax*.

```
QUADINSTR <- { 'quad' } S QUAD
S <- %s+
QUAD <- { [0-9a-fA-F]^8 }
```

A.23.5. *Bytecode Encoding*. This instruction has a size of 8 bytes and is structured in bytecode as follows:

00	00	04	03	7e
04	00	00	00	00

A.23.6. *Execution State Change*. .

A.24. Instruction: **range**.

A.24.1. *Mode*. This instruction is available in mode 0 (parser).

A.24.2. *Summary*.

A.24.3. *Grammar and Compiling*.

A.24.4. *Assembly Syntax*.

```
RANGEINSTR <- { 'range' } S UNSIGNED S UNSIGNED
S <- %s+
UNSIGNED <- { [0-9]+ }
```

A.24.5. *Bytecode Encoding*. This instruction has a size of 12 bytes and is structured in bytecode as follows:

00	00	08	03	bd
04	00	00	00	00
08	00	00	00	00

A.24.6. *Execution State Change*. .

A.25. Instruction: **replace**.

A.25.1. *Mode*. This instruction is available in mode 0 (parser).

A.25.2. *Summary*.

A.25.3. *Grammar and Compiling*.

A.25.4. *Assembly Syntax*.

```
REPLACEINSTR <- { 'replace' } S LABEL S LABEL
S <- %s+
LABEL <- { [a-zA-Z0-9_]^1-64 }
```

A.25.5. *Bytecode Encoding*. This instruction has a size of 12 bytes and is structured in bytecode as follows:

00	00	08	03	48
04	00	00	00	00
08	00	00	00	00

A.25.6. *Execution State Change*. .

A.26. **Instruction: ret.**

A.26.1. *Mode.* This instruction is available in mode 0 (parser).

A.26.2. *Summary.*

A.26.3. *Grammar and Compiling.*

A.26.4. *Assembly Syntax.*

```
RETINSTR <- { 'ret' }
```

A.26.5. *Bytecode Encoding.* This instruction has a size of 4 bytes and is structured in bytecode as follows:

00	00	00	03	a0
----	----	----	----	----

A.26.6. *Execution State Change.* .

A.27. **Instruction: `scr_add`.**

A.27.1. *Mode*. This instruction is available in mode 1 (script interpreter).

A.27.2. *Summary*.

A.27.3. *Grammar and Compiling*.

A.27.4. *Assembly Syntax*.

```
SCR_ADD <- { 'scr_add' }
```

A.27.5. *Bytecode Encoding*. This instruction has a size of 4 bytes and is structured in bytecode as follows:

00	00	00	05	0c
----	----	----	----	----

A.27.6. *Execution State Change*. .

A.28. Instruction: `scr_array`.

A.28.1. *Mode*. This instruction is available in mode 1 (script interpreter).

A.28.2. *Summary*.

A.28.3. *Grammar and Compiling*.

A.28.4. *Assembly Syntax*.

```
SCR_ARRAY <- { 'scr_array' } S UNSIGNED
S <- %s+
UNSIGNED <- { [0-9]+ }
```

A.28.5. *Bytecode Encoding*. This instruction has a size of 8 bytes and is structured in bytecode as follows:

00	00	04	00	06
04	00	00	00	00

A.28.6. *Execution State Change*. .

A.29. Instruction: `scr_assign`.

A.29.1. *Mode*. This instruction is available in mode 1 (script interpreter).

A.29.2. *Summary*.

A.29.3. *Grammar and Compiling*.

A.29.4. *Assembly Syntax*.

```
SCR_ASSIGN <- { 'scr_assign' }
```

A.29.5. *Bytecode Encoding*. This instruction has a size of 4 bytes and is structured in bytecode as follows:

00 00 00 05 c9

A.29.6. *Execution State Change*. .

A.30. **Instruction: `scr_bitand`.**

A.30.1. *Mode.* This instruction is available in mode 1 (script interpreter).

A.30.2. *Summary.*

A.30.3. *Grammar and Compiling.*

A.30.4. *Assembly Syntax.*



A.30.5. *Bytecode Encoding.* This instruction has a size of 4 bytes and is structured in bytecode as follows:

00	00	00	05	27
----	----	----	----	----

A.30.6. *Execution State Change.* .

A.31. **Instruction: `scr_bitnot`.**

A.31.1. *Mode*. This instruction is available in mode 1 (script interpreter).

A.31.2. *Summary*.

A.31.3. *Grammar and Compiling*.

A.31.4. *Assembly Syntax*.

A.31.5. *Bytecode Encoding*. This instruction has a size of 4 bytes and is structured in bytecode as follows:

00	00	00	05	74
----	----	----	----	----

A.31.6. *Execution State Change*. .

A.32. Instruction: `scr_bitor`.

A.32.1. *Mode*. This instruction is available in mode 1 (script interpreter).

A.32.2. *Summary*.

A.32.3. *Grammar and Compiling*.

A.32.4. *Assembly Syntax*.



A.32.5. *Bytecode Encoding*. This instruction has a size of 4 bytes and is structured in bytecode as follows:

00 00 00 05 3c

A.32.6. *Execution State Change*. .

A.33. **Instruction: `scr_bitxor`.**

A.33.1. *Mode*. This instruction is available in mode 1 (script interpreter).

A.33.2. *Summary*.

A.33.3. *Grammar and Compiling*.

A.33.4. *Assembly Syntax*.

A.33.5. *Bytecode Encoding*. This instruction has a size of 4 bytes and is structured in bytecode as follows:

00	00	00	05	2d
----	----	----	----	----

A.33.6. *Execution State Change*. .

A.34. Instruction: `scr_builtin`.

A.34.1. *Mode*. This instruction is available in mode 1 (script interpreter).

A.34.2. *Summary*.

A.34.3. *Grammar and Compiling*.

A.34.4. *Assembly Syntax*.

```
SCR_BUILTIN <- { 'scr_builtin' } S UNSIGNED
S <- %s+
UNSIGNED <- { [0-9]+ }
```

A.34.5. *Bytecode Encoding*. This instruction has a size of 8 bytes and is structured in bytecode as follows:

00	00	04	07	cf
04	00	00	00	00

A.34.6. *Execution State Change*. .

A.35. Instruction: `scr_call`.

A.35.1. *Mode*. This instruction is available in mode 1 (script interpreter).

A.35.2. *Summary*.

A.35.3. *Grammar and Compiling*.

A.35.4. *Assembly Syntax*.

```
SCR_CALL <- { 'scr_call' } S LABEL
S <- %s+
LABEL <- { [a-zA-Z0-9_]^1-64 }
```

A.35.5. *Bytecode Encoding*. This instruction has a size of 8 bytes and is structured in bytecode as follows:

00	00	04	05	03
04	00	00	00	00

A.35.6. *Execution State Change*. .

A.36. Instruction: `scr_condjump`.

A.36.1. *Mode*. This instruction is available in mode 1 (script interpreter).

A.36.2. *Summary*.

A.36.3. *Grammar and Compiling*.

A.36.4. *Assembly Syntax*.

```
SCR_CONDJUMP <- { 'scr_condjump' } S LABEL
S <- %s+
LABEL <- { [a-zA-Z0-9_]^1-64 }
```

A.36.5. *Bytecode Encoding*. This instruction has a size of 8 bytes and is structured in bytecode as follows:

00	00	04	00	0c
04	00	00	00	00

A.36.6. *Execution State Change*. .

A.37. **Instruction: `scr_dec`.**

A.37.1. *Mode*. This instruction is available in mode 1 (script interpreter).

A.37.2. *Summary*.

A.37.3. *Grammar and Compiling*.

A.37.4. *Assembly Syntax*.



A.37.5. *Bytecode Encoding*. This instruction has a size of 4 bytes and is structured in bytecode as follows:

00	00	00	05	3a
----	----	----	----	----

A.37.6. *Execution State Change*. .

A.38. Instruction: `scr_div`.

A.38.1. *Mode*. This instruction is available in mode 1 (script interpreter).

A.38.2. *Summary*.

A.38.3. *Grammar and Compiling*.

A.38.4. *Assembly Syntax*.



A.38.5. *Bytecode Encoding*. This instruction has a size of 4 bytes and is structured in bytecode as follows:

00

00

00

05

81

A.38.6. *Execution State Change*. .

A.39. Instruction: `scr_equals`.

A.39.1. *Mode*. This instruction is available in mode 1 (script interpreter).

A.39.2. *Summary*.

A.39.3. *Grammar and Compiling*.

A.39.4. *Assembly Syntax*.

```
SCR_EQUALS <- { 'scr_equals' }
```

A.39.5. *Bytecode Encoding*. This instruction has a size of 4 bytes and is structured in bytecode as follows:

00 00 00 6c

A.39.6. *Execution State Change*. .

A.40. Instruction: `scr.gt`.

A.40.1. *Mode*. This instruction is available in mode 1 (script interpreter).

A.40.2. *Summary*.

A.40.3. *Grammar and Compiling*.

A.40.4. *Assembly Syntax*.



A.40.5. *Bytecode Encoding*. This instruction has a size of 4 bytes and is structured in bytecode as follows:

00 00 00 05 6f

A.40.6. *Execution State Change*. .

A.41. **Instruction: `scr_gteq`.**

A.41.1. *Mode*. This instruction is available in mode 1 (script interpreter).

A.41.2. *Summary*.

A.41.3. *Grammar and Compiling*.

A.41.4. *Assembly Syntax*.



A.41.5. *Bytecode Encoding*. This instruction has a size of 4 bytes and is structured in bytecode as follows:

00	00	00	05	4d
----	----	----	----	----

A.41.6. *Execution State Change*. .

A.42. Instruction: `scr_inc`.

A.42.1. *Mode*. This instruction is available in mode 1 (script interpreter).

A.42.2. *Summary*.

A.42.3. *Grammar and Compiling*.

A.42.4. *Assembly Syntax*.



A.42.5. *Bytecode Encoding*. This instruction has a size of 4 bytes and is structured in bytecode as follows:

00

00

00

05

f6

A.42.6. *Execution State Change*. .

A.43. **Instruction: scr_index.**

A.43.1. *Summary.*

A.43.2. *Grammar and Compiling.*

A.43.3. *Assembly Syntax.*

```
SCR_INDEX <- { 'scr_index' }
```

A.43.4. *Bytecode Encoding.* This instruction has a size of 4 bytes and is structured in bytecode as follows:

00	00	00	09
----	----	----	----

A.43.5. *Execution State Change.* .

A.44. Instruction: `scr_logand`.

A.44.1. *Mode*. This instruction is available in mode 1 (script interpreter).

A.44.2. *Summary*.

A.44.3. *Grammar and Compiling*.

A.44.4. *Assembly Syntax*.



A.44.5. *Bytecode Encoding*. This instruction has a size of 4 bytes and is structured in bytecode as follows:

00	00	00	05	2e
----	----	----	----	----

A.44.6. *Execution State Change*. .

A.45. **Instruction:** `scr_lognot`.

A.45.1. *Mode*. This instruction is available in mode 1 (script interpreter).

A.45.2. *Summary*.

A.45.3. *Grammar and Compiling*.

A.45.4. *Assembly Syntax*.



A.45.5. *Bytecode Encoding*. This instruction has a size of 4 bytes and is structured in bytecode as follows:

00	00	00	05	f9
----	----	----	----	----

A.45.6. *Execution State Change*. .

A.46. Instruction: `scr_logor`.

A.46.1. *Mode*. This instruction is available in mode 1 (script interpreter).

A.46.2. *Summary*.

A.46.3. *Grammar and Compiling*.

A.46.4. *Assembly Syntax*.



A.46.5. *Bytecode Encoding*. This instruction has a size of 4 bytes and is structured in bytecode as follows:

00	00	00	05	a9
----	----	----	----	----

A.46.6. *Execution State Change*. .

A.47. **Instruction: scr_lt.**

A.47.1. *Mode.* This instruction is available in mode 1 (script interpreter).

A.47.2. *Summary.*

A.47.3. *Grammar and Compiling.*

A.47.4. *Assembly Syntax.*

```
SCR_LT <- { 'scr_lt' }
```

A.47.5. *Bytecode Encoding.* This instruction has a size of 4 bytes and is structured in bytecode as follows:

00	00	00	05	95
----	----	----	----	----

A.47.6. *Execution State Change.* .

A.48. Instruction: `scr_lteq`.

A.48.1. *Mode*. This instruction is available in mode 1 (script interpreter).

A.48.2. *Summary*.

A.48.3. *Grammar and Compiling*.

A.48.4. *Assembly Syntax*.



A.48.5. *Bytecode Encoding*. This instruction has a size of 4 bytes and is structured in bytecode as follows:

00 00 00 05 22

A.48.6. *Execution State Change*. .

A.49. **Instruction: scr_mul.**

A.49.1. *Mode.* This instruction is available in mode 1 (script interpreter).

A.49.2. *Summary.*

A.49.3. *Grammar and Compiling.*

A.49.4. *Assembly Syntax.*

```
SCR_MUL <- { 'scr_mul' }
```

A.49.5. *Bytecode Encoding.* This instruction has a size of 4 bytes and is structured in bytecode as follows:

00	00	00	05	8b
----	----	----	----	----

A.49.6. *Execution State Change.* .

A.50. Instruction: `scr_nequals`.

A.50.1. *Mode*. This instruction is available in mode 1 (script interpreter).

A.50.2. *Summary*.

A.50.3. *Grammar and Compiling*.

A.50.4. *Assembly Syntax*.



A.50.5. *Bytecode Encoding*. This instruction has a size of 4 bytes and is structured in bytecode as follows:

00	00	00	05	72
----	----	----	----	----

A.50.6. *Execution State Change*. .

A.51. **Instruction: scr_pop.**

A.51.1. *Mode.* This instruction is available in mode 1 (script interpreter).

A.51.2. *Summary.*

A.51.3. *Grammar and Compiling.*

A.51.4. *Assembly Syntax.*

```
SCR_POP <- { 'scr_pop' }
```

A.51.5. *Bytecode Encoding.* This instruction has a size of 4 bytes and is structured in bytecode as follows:

00

00

00

05

cc

A.51.6. *Execution State Change.* .

A.52. Instruction: `scr_pow`.

A.52.1. *Mode*. This instruction is available in mode 1 (script interpreter).

A.52.2. *Summary*.

A.52.3. *Grammar and Compiling*.

A.52.4. *Assembly Syntax*.



A.52.5. *Bytecode Encoding*. This instruction has a size of 4 bytes and is structured in bytecode as follows:

00

00

00

05

42

A.52.6. *Execution State Change*. .

A.53. Instruction: `scr_push`.

A.53.1. *Mode*. This instruction is available in mode 1 (script interpreter).

A.53.2. *Summary*.

A.53.3. *Grammar and Compiling*.

A.53.4. *Assembly Syntax*.

```
SCR_PUSH <- { 'scr_push' } S LITERAL
S <- %s+
LITERAL <- REGISTERREF /
```

A.53.5. *Bytecode Encoding*. This instruction has a size of 16 bytes and is structured in bytecode as follows:

00	00	0c	00	03
04	00	00	00	00
08	00	00	00	00
12	00	00	00	00

A.53.6. *Execution State Change*. .

A.54. Instruction: `scr_ret`.

A.54.1. *Mode*. This instruction is available in mode 1 (script interpreter).

A.54.2. *Summary*.

A.54.3. *Grammar and Compiling*.

A.54.4. *Assembly Syntax*.

```
SCR_RET <- { 'scr_ret' }
```

A.54.5. *Bytecode Encoding*. This instruction has a size of 4 bytes and is structured in bytecode as follows:

00	00	00	05	55
----	----	----	----	----

A.54.6. *Execution State Change*. .

A.55. **Instruction: scr_shift.**

A.55.1. *Mode.* This instruction is available in mode 1 (script interpreter).

A.55.2. *Summary.*

A.55.3. *Grammar and Compiling.*

A.55.4. *Assembly Syntax.*

```
SCR_SHIFT <- { 'scr_shift' } S UNSIGNED
S <- %s+
UNSIGNED <- { [0-9]+ }
```

A.55.5. *Bytecode Encoding.* This instruction has a size of 8 bytes and is structured in bytecode as follows:

00	00	04	00	05
04	00	00	00	00

A.55.6. *Execution State Change.* .

A.56. Instruction: `scr_shiftin`.

A.56.1. *Mode*. This instruction is available in mode 1 (script interpreter).

A.56.2. *Summary*.

A.56.3. *Grammar and Compiling*.

A.56.4. *Assembly Syntax*.



A.56.5. *Bytecode Encoding*. This instruction has a size of 4 bytes and is structured in bytecode as follows:

00	00	00	05	7d
----	----	----	----	----

A.56.6. *Execution State Change*. .

A.57. **Instruction: `scr_shiftout`.**

A.57.1. *Mode*. This instruction is available in mode 1 (script interpreter).

A.57.2. *Summary*.

A.57.3. *Grammar and Compiling*.

A.57.4. *Assembly Syntax*.



A.57.5. *Bytecode Encoding*. This instruction has a size of 4 bytes and is structured in bytecode as follows:

00	00	00	05	17
----	----	----	----	----

A.57.6. *Execution State Change*. .

A.58. Instruction: `scr_string`.

A.58.1. *Mode*. This instruction is available in mode 1 (script interpreter).

A.58.2. *Summary*.

A.58.3. *Grammar and Compiling*.

A.58.4. *Assembly Syntax*.

```
SCR_STRING <- { 'scr_string' } S STRINGLITERAL
S <- %s+
STRINGLITERAL <- '\'' { ( '\\' ([nrtv\\'] / [0-9]^3) / [^'\\] )* } '\'
```

A.58.5. *Bytecode Encoding*. This instruction has a size of 4 bytes and is structured in bytecode as follows:

00	00	00	17	bb
----	----	----	----	----

A.58.6. *Execution State Change*. .

A.59. **Instruction: `scr_sub`.**

A.59.1. *Mode*. This instruction is available in mode 1 (script interpreter).

A.59.2. *Summary*.

A.59.3. *Grammar and Compiling*.

A.59.4. *Assembly Syntax*.

A.59.5. *Bytecode Encoding*. This instruction has a size of 4 bytes and is structured in bytecode as follows:

00	00	00	05	bb
----	----	----	----	----

A.59.6. *Execution State Change*. .

A.60. Instruction: set.

A.60.1. *Mode.* This instruction is available in mode 0 (parser).

A.60.2. *Summary.*

A.60.3. *Grammar and Compiling.*

A.60.4. *Assembly Syntax.*

```
SETINSTR <- { 'set' } S SET
S <- %s+
SET <- { [0-9a-fA-F]^64 }
```

A.60.5. *Bytecode Encoding.* This instruction has a size of 36 bytes and is structured in bytecode as follows:

00	00	20	03	ca
04	00	00	00	00
08	00	00	00	00
12	00	00	00	00
16	00	00	00	00
20	00	00	00	00
24	00	00	00	00
28	00	00	00	00
32	00	00	00	00

A.60.6. *Execution State Change.* .

A.61. Instruction: skip.

A.61.1. *Mode*. This instruction is available in mode 0 (parser).

A.61.2. *Summary*.

A.61.3. *Grammar and Compiling*.

A.61.4. *Assembly Syntax*.

```
SKIPINSTR <- { 'skip' } S UNSIGNED
S <- %s+
UNSIGNED <- { [0-9]+ }
```

A.61.5. *Bytecode Encoding*. This instruction has a size of 8 bytes and is structured in bytecode as follows:

00	00	04	03	30
04	00	00	00	00

A.61.6. *Execution State Change*. .

A.62. Instruction: **span**.

A.62.1. *Mode*. This instruction is available in mode 0 (parser).

A.62.2. *Summary*.

A.62.3. *Grammar and Compiling*.

A.62.4. *Assembly Syntax*.

```
SPANINSTR <- { 'span' } S SET
S <- %s+
SET <- { [0-9a-fA-F]^64 }
```

A.62.5. *Bytecode Encoding*. This instruction has a size of 36 bytes and is structured in bytecode as follows:

00	00	20	03	e1
04	00	00	00	00
08	00	00	00	00
12	00	00	00	00
16	00	00	00	00
20	00	00	00	00
24	00	00	00	00
28	00	00	00	00
32	00	00	00	00

A.62.6. *Execution State Change*. .

A.63. **Instruction: testany.**

A.63.1. *Mode.* This instruction is available in mode 0 (parser).

A.63.2. *Summary.*

A.63.3. *Grammar and Compiling.*

A.63.4. *Assembly Syntax.*

```
TESTANYINSTR <- { 'testany' } S LABEL
S <- %s+
LABEL <- { [a-zA-Z0-9_]^1-64 }
```

A.63.5. *Bytecode Encoding.* This instruction has a size of 8 bytes and is structured in bytecode as follows:

00	00	04	03	06
04	00	00	00	00

A.63.6. *Execution State Change.* .

A.64. Instruction: **testchar**.

A.64.1. *Mode*. This instruction is available in mode 0 (parser).

A.64.2. *Summary*.

A.64.3. *Grammar and Compiling*.

A.64.4. *Assembly Syntax*.

```
TESTCHARINSTR <- { 'testchar' } S HEXBYTE S LABEL
S <- %s+
HEXBYTE <- { [0-9a-fA-F]^2 }
LABEL <- { [a-zA-Z0-9_]^1-64 }
```

A.64.5. *Bytecode Encoding*. This instruction has a size of 12 bytes and is structured in bytecode as follows:

00	00	08	03	9a
04	00	00	00	00
08	00	00	00	00

A.64.6. *Execution State Change*. .

A.65. Instruction: testquad.

A.65.1. *Mode*. This instruction is available in mode 0 (parser).

A.65.2. *Summary*.

A.65.3. *Grammar and Compiling*.

A.65.4. *Assembly Syntax*.

```
TESTQUADINSTR <- { 'testquad' } S QUAD S LABEL
S <- %s+
QUAD <- { [0-9a-fA-F]^8 }
LABEL <- { [a-zA-Z0-9_]^1-64 }
```

A.65.5. *Bytecode Encoding*. This instruction has a size of 12 bytes and is structured in bytecode as follows:

00	00	08	03	db
04	00	00	00	00
08	00	00	00	00

A.65.6. *Execution State Change*. .

A.66. Instruction: testset.

A.66.1. *Mode*. This instruction is available in mode 0 (parser).

A.66.2. *Summary*.

A.66.3. *Grammar and Compiling*.

A.66.4. *Assembly Syntax*.

```
TESTSETINSTR <- { 'testset' } S SET S LABEL
S <- %s+
SET <- { [0-9a-fA-F]^64 }
LABEL <- { [a-zA-Z0-9_]^1-64 }
```

A.66.5. *Bytecode Encoding*. This instruction has a size of 40 bytes and is structured in bytecode as follows:

00	00	24	03	63
04	00	00	00	00
08	00	00	00	00
12	00	00	00	00
16	00	00	00	00
20	00	00	00	00
24	00	00	00	00
28	00	00	00	00
32	00	00	00	00
36	00	00	00	00

A.66.6. *Execution State Change*. .

A.67. **Instruction: trap.**

A.67.1. *Mode.* This instruction is available in both mode 0 (parser) and mode 1 (script interpreter).

A.67.2. *Summary.*

A.67.3. *Grammar and Compiling.*

A.67.4. *Assembly Syntax.*

```
TRAPINSTR <- { 'trap' }
```

A.67.5. *Bytecode Encoding.* This instruction has a size of 4 bytes and is structured in bytecode as follows:

00	ff	00	ff
----	----	----	----

A.67.6. *Execution State Change.* .

A.68. **Instruction: var.**

A.68.1. *Mode.* This instruction is available in mode 0 (parser).

A.68.2. *Summary.*

A.68.3. *Grammar and Compiling.*

A.68.4. *Assembly Syntax.*

```
VARINSTR <- { 'var' } S SLOT
S <- %s+
SLOT <- UNSIGNED
```

A.68.5. *Bytecode Encoding.* This instruction has a size of 8 bytes and is structured in bytecode as follows:

00	00	04	03	ee
04	00	00	00	00

A.68.6. *Execution State Change.* .